Reproducing and Improving QUANT: A Minimalist Interval Method for Time Series Classification

Dániel Sajben
University College Dublin
Dublin, Ireland
Student ID: 26107621
daniel.sajben@ucdconnect.ie

Tymur Sukhariev
University College Dublin
Dublin, Ireland
Student ID: 26100566
tymur.sukhariev@ucdconnect.ie

***Abstract*—QUANT is a time series classifier that describes each series by the quantiles of a fixed set of windows and passes them to an off-the-shelf tree ensemble; its authors report state-of-the-art accuracy for interval methods on 142 UCR datasets in under 15 minutes on one CPU core. Re-running the published code on the published data, we got 85.51% mean accuracy against the paper's 85.68%; the gap comes from four datasets corrected after publication. Controlled tests showed that QUANT's fixed windows make it sensitive to patterns that appear at different times, and that it falls behind on datasets with many classes and few examples. Of the fifteen ideas we tested to fix this, the best lets a second, position-free method (Hydra) vote only when QUANT is unsure: 87.16% (+1.8 points, $p < 0.001$) at 2.1× QUANT's compute time, level with MultiRocket+Hydra at less than half its reported cost and 0.4 points below HIVE-COTE 2, which needs weeks. Adding two of our feature ideas on top gives 87.47% (not timed), statistically level with HIVE-COTE 2. On 19 multivariate datasets, new views that combine channels raise QUANT's accuracy by 2.6 points (3.6 with Hydra) and make it far less sensitive to tilted sensors. Code, commands and results: https://github.com/SajbenDani/AI-For-Timeseries-Project.**

***Index Terms*—time series classification, interval methods, quantiles, reproducibility, UCR archive, multivariate time series**

I. Introduction

The problem. In time series classification (TSC), each example is a sequence of measurements over time, such as an ECG trace or an accelerometer recording, and the goal is to predict its label. The most accurate TSC methods, such as HIVE-COTE 2 (HC2) [1], take days or weeks of compute on the standard UCR benchmark [2]. This makes them hard to use in practice.

A. Motivation

In a world where the trend is to build bigger and bigger data centres to gain one or two percentage points on a benchmark, QUANT [3], an efficiency focused method, really stood out to us. We see the same habit in how people use language models. Many pick the largest model available (for example Claude Fable over Opus or Sonnet) even for simple tasks, because it is supposed to be the best and they trust it most. The strain on the environment and the size of the bill rarely come up, at least as long as someone else is paying. We think efficiency is key for the future, and we expect the field to keep moving in that direction. Methods that use only the compute a task actually needs, from mixture-of-experts models to small, specialised methods like QUANT, are a good step that way.

Unsurprisingly, QUANT's main strength is efficiency. Instead of adding more parts, the authors stepped back and asked what an interval method actually needs. Their answer was very little: one type of feature (quantiles), a fixed set of windows and an off-the-shelf classifier. What sold the paper to us was the runtime comparison. QUANT needs under 15 minutes on a single CPU core for all 142 datasets. HC2 [1], the most accurate method in the comparison, needs several weeks and scores only about 2 percentage points higher (0.879 vs 0.857). That is exactly the trade-off we see with large language models. We wanted to find out whether we could shrink this gap further, or make QUANT even cheaper without losing accuracy.

The paper also compares QUANT with several other methods, such as MultiRocket+Hydra, rSTSF [4], RDST and HC2. All of them are slower, from about an hour for MultiRocket+Hydra and rSTSF to several weeks for HC2, and the most accurate ones beat QUANT by up to about 2 points. We thought we could take ideas from these methods and build them into QUANT, raising its accuracy while keeping it fast. We also liked that every method is tested on exactly the same 142 datasets and 30 resamples. The comparisons felt transparent, so we believed we had a real chance to improve QUANT in a statistically significant way. That could even help the time series community, which mattered more to us than tweaking random seeds for an extra 0.01% just to finish an assignment.

In short, this is why we chose this paper. We hope you are as curious about our results as we were when we wrote this section, before we started the research.

B. How QUANT Works

QUANT [3] is a classification model: it learns from labelled example series and then predicts the class of new ones. It

works in four steps.

1) **Four views.** It looks at every series in four ways: the raw values; the first difference, i.e. how much the series changes from one step to the next (lightly smoothed); the second difference; and the magnitudes of its Fourier transform, i.e. which frequencies it contains.
2) **Fixed windows.** Each view is cut into 1, 2, 4, 8, 16 and 32 equal windows, plus the same windows shifted by half a window.
3) **Quantiles.** In each window QUANT records about one evenly spaced quantile per four values (the median is the 50% quantile, the maximum the 100% one). Every second quantile is taken relative to the window mean, so the features describe both the level and the shape of the values.
4) **Classifier.** A forest of 200 extremely randomised trees (ExtraTrees [5]) learns which of these numbers separate the classes.

For a heartbeat recording, one feature might be "the 90% quantile of the raw values in the second quarter", which is high when a tall beat falls there. Long windows describe *which* values occur, short ones *where*. The transform learns nothing; the forest does all the learning and takes most of the time. QUANT is part of the aeon toolkit [6] and the most accurate method on the large MONSTER datasets [7].

C. The Benchmark

We use the paper's 142 univariate datasets from the "bake off redux" [2]: the 112 equal-length datasets of the 2018 UCR archive [8] and 30 newer ones, from outlines and sensors to ECG, spectra and astronomy (2 to 60 classes, 16 to 8,926 training series). Each comes with 30 fixed train/test splits ("resamples"), the original one plus 29 with the same class balance, so methods can be compared dataset by dataset. Sixteen new "Eq" datasets were made equal-length, mostly by padding: 11 are our "padded" datasets, three (Asphalt) our unseen data. Section VI adds 19 multivariate datasets.

D. Where QUANT Is Weak

Section IV shows that QUANT suffers when the same pattern appears at different positions (-5.6 percentage points with series placed at random positions, Hydra [9] only -2.4) and with many classes and few examples (5.6 points behind Hydra).

E. Contributions

(1) A scripted reproduction on a laptop CPU, listing what is missing from the paper or has changed since. (2) Controlled tests of where QUANT is weak. (3) Fifteen improvement ideas tested with the paper's protocol, including a gated QUANT+Hydra ensemble with top-method accuracy at $2.1\times$ QUANT's cost. (4) Cross-channel views for multivariate data, tested for accuracy and on tilted sensors. (5) Negative results and checks on unseen data, reported as they came out.

TABLE I
MEAN ACCURACY OF QUANT OVER 30 RESAMPLES

Datasets	n	Paper	Bake-off	Ours
Original UCR	112	0.8671	0.8671	0.8671
New, unchanged	26	0.8008	0.8019	0.8017
New, corrected later	4	0.9306	0.8662	0.8673
All	142	0.8568	0.8552	0.8551
All except corrected	138	0.8546	0.8548	0.8548

Paper: published per-resample results. Bake-off: independent QUANT run by [2] on the corrected data.

II. Related Methods

ROCKET [10], MiniRocket [11] and MultiRocket [12] slide thousands of random convolution kernels (small pattern detectors) along each series, summarise every output over the whole series, so it does not matter *where* a pattern occurs, and train a linear (ridge) classifier. Hydra [9], from QUANT's authors, counts which kernel in each group matches best at every time point, again ignoring position; MultiRocket+Hydra [9] came second in the bake-off redux. The winner, HC2 [1], combines four very different classifiers weighted by their estimated accuracy [13]. rSTSF [4] and DrCIF [1] are interval methods like QUANT.

III. Reproducing QUANT

A. Approach

We ran the authors' code [14] unchanged (commit `f3907f7`) through a thin wrapper that also records timings, saves after every resample and fixes the forest's random seed (Python 3.11, PyTorch 2.14 on CPU, scikit-learn 1.9.1). A dataset's accuracy is the mean over its 30 resamples. We compare with the paper's published per-resample results and with an independent QUANT run by the bake-off authors [2], who also provide the competitors' results. A Welch t-test over the resamples tells whether a dataset's gap exceeds resampling noise; across datasets we report the mean, wins/draws/losses (W/D/L) and a Wilcoxon signed-rank test, as the paper does [15] (a small p means the difference is unlikely to be luck).

B. Results

Our reproduction of QUANT. We got 85.51% against the paper's 85.68% (Table I). On the 112 original datasets our mean is 86.71%, the same as the paper's to four decimals; no dataset differs by more than resampling noise (all Welch $p > 0.1$), and our per-resample accuracies track the paper's (median correlation 0.89). The gap comes from Covid3Month and the three FloodModeling sets, which were corrected after the QUANT preprint appeared; only the new versions can be downloaded, and on them QUANT scores much lower (Covid3Month: 0.62 vs 0.81). The bake-off authors' own QUANT run on the corrected data matches ours, so the gap comes from the data, not from the method or our code.

Against other methods (Table II), every conclusion of the paper still holds: QUANT is slightly ahead of the best interval

TABLE II
QUANT VS OTHER METHODS (BAKE-OFF RESULTS): MEAN ACCURACY DIFFERENCE (Δ, POSITIVE = QUANT BETTER), WINS/DRAWS/LOSSES AND WILCOXON p

Method	n	Δ paper	Δ ours	W/D/L ours	p ours
HC2	142	-0.0222	-0.0207	42/7/93	<0.001
MultiRocket+Hydra	142	-0.0165	-0.0163	48/6/88	<0.001
RDST [16]	142	-0.0083	-0.0077	62/4/76	0.133
WEASEL-D [17]	142	-0.0040	-0.0022	69/3/70	0.493
InceptionTime [18]	142	$+0.0001$	-0.0014	64/4/74	0.252
FreshPRINCE [19]	142	$+0.0087$	$+0.0087$	99/3/40	<0.001
Proximity Forest [20]	142	$+0.0267$	$+0.0270$	97/6/39	<0.001
rSTSF	112	$+0.0031$	$+0.0031$	66/8/38	0.022
DrCIF	112	$+0.0035$	$+0.0034$	60/8/44	0.157

methods (significantly only against rSTSF), roughly level with RDST, WEASEL-D and InceptionTime, and clearly behind HC2 and MultiRocket+Hydra.

Runtime. One pass over all 142 datasets takes 18.3 minutes on one core of our laptop CPU (AMD Ryzen 5 5600H), against 14.9 minutes on server CPUs in the paper. About 78% of the time goes into training the forest.

C. Lessons Learned

The accuracy claims, the code and the ranking all hold, and reproducing them needs only a laptop CPU, but some details are missing or have changed. (1) **The resamples are not documented:** of the two sets of split indices (from the Java and the Python bake-off tool), per-resample correlations show that the paper used the Java set for the 112 original datasets (0.83–0.96, against about 0 for Python) and the Python set for the 30 new ones. (2) **Dataset names differ** between the code and the download, so 24 datasets fail without renaming. (3) **The data changed after publication**, so benchmark numbers should come with a data version. (4) **The forest is not seeded**: results change between runs (standard deviation up to 0.09 on single splits of Beef). Even with a fixed seed they differ between computers: on our second laptop 722 of 1,420 QUANT splits changed, by up to 4.3 points, with the same mean. (5) **W/D/L counts depend on rounding** (1–3 datasets move). (6) **The sensitivity analysis cannot be re-run**: its 50 datasets and folds are not listed. **Takeaway.** The published results hold, but gaps between strong methods are small and noisy, so any improvement has to be tested over many resamples with paired tests.

IV. WHERE QUANT IS WEAK

To decide what to fix, we changed one thing at a time and watched QUANT's accuracy, with Hydra [9] as a reference (Fig. 1).

Patterns at different positions. We placed every series of 30 datasets at a random position inside a window 10%, 25% or 50% longer, filling the gap with the series' first or last value plus a little noise. At 50%, QUANT loses 5.6 points and Hydra 2.4 (Fig. 1a), with the worst cases on face outlines and handwritten words (FaceFour -22, WordSynonyms -20). The cause is the fixed windows: a feature such as "the 90th percentile between 25% and 50% of the series" changes when the pattern moves. On the 10 gesture datasets, where people start moving at different times, QUANT trails MultiRocket+Hydra by 5.1 points, against 1.7 elsewhere, as Zhang et al. [21] found for an older interval method.

Many classes and few examples. Cutting 25 datasets down to 2–20 training examples per class barely changes QUANT's gap to Hydra (0.2 to 1.1 points; Fig. 1b), and a linear classifier on QUANT's features does even worse. But on the 18 datasets with many classes *and* few examples per class, QUANT trails Hydra by 5.6 points (Hydra wins 15 of 18, $p = 0.002$), and on the Pig recordings (52 classes, 2 examples each) by 8–27 points. We think the cause is position again: with two examples per class, the forest cannot learn every place where a heartbeat may sit.

Padding. On the 11 padded datasets, the more heavily padded test series are about 8 points harder for *both* methods, and QUANT's gap to Hydra hardly changes (Fig. 1c). Padding hurts, but it is not what sets QUANT apart.

Other findings. (i) Removing test series that copy training series costs QUANT up to 5.7 points on one dataset (Worms) but only 0.09 on the mean. (ii) Normalising each series first costs 3.1 points, because the absolute level often *is* the class.

V. IMPROVING QUANT

A. How We Tested

Data and splits. The same 142 datasets and official resamples, but only the first 10 resamples, to fit the runs on a laptop; plain QUANT scores 85.36% on these (85.51% on all 30). The forest's seed is the resample number, so our code reproduces Section III's QUANT accuracies exactly on the same computer. Ideas 7, 9, 11–13 and the stacking ran on our second laptop and are compared with plain QUANT fitted there. We never change QUANT: each idea *adds* features, training copies or a second model, except ideas 9 and 17, which drop or rescale features, and idea 13, which swaps the classifier.

Statistics. Per dataset we average the 10 resamples; across datasets we report W/D/L (four decimals) and a Wilcoxon signed-rank test, with $p < 0.05$ as a real effect. Every gain in Table III's "Worked" group stays significant after a Holm correction (a stricter test for trying many ideas) over all 21 comparisons in the table, ideas 2, 10 and 14 only just (corrected $p \approx 0.049$); the test-time-averaging loss does not. Inside the training data, idea 3 uses cross-validation (up to 5 folds) and idea 18 leave-one-out predictions.

Stress tests. On 30 datasets $\times$ 5 resamples: the random-position test of Section IV, and a speed test in which every series is played faster at the start and slower at the end, or the reverse (time $t \mapsto t^g$, $|\log g| \leq 0.6$).

Cost is one-core compute time (same CPU) for features, training and prediction on one split of every dataset, relative to plain QUANT in the same run (about 16 minutes; its time varied from 15.7 to 18.3 minutes between runs). The two

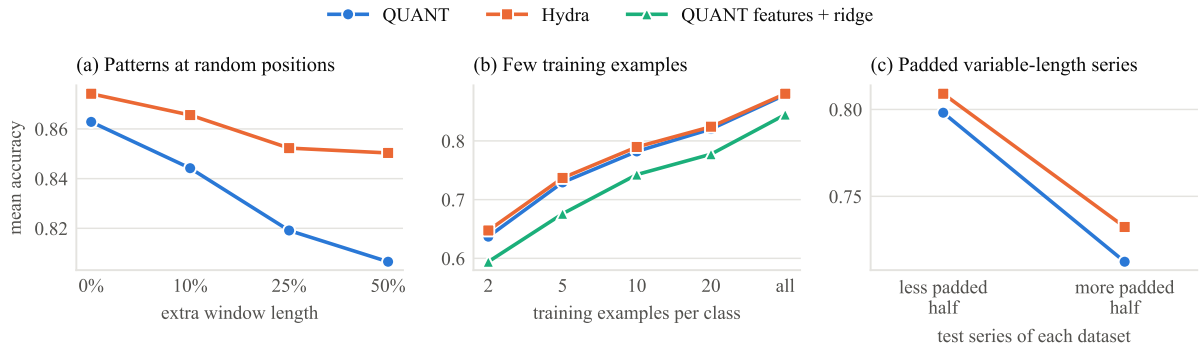


Fig. 1. Where QUANT is weak, with Hydra as a reference. (a) Every series is placed at a random position inside a longer window (30 datasets, 10 resamples). (b) Training sets cut to a few examples per class (25 datasets); "QUANT features + ridge" replaces the forest by a linear classifier. (c) The 11 padded variable-length datasets: accuracy on the less and the more padded half of each dataset's test series.

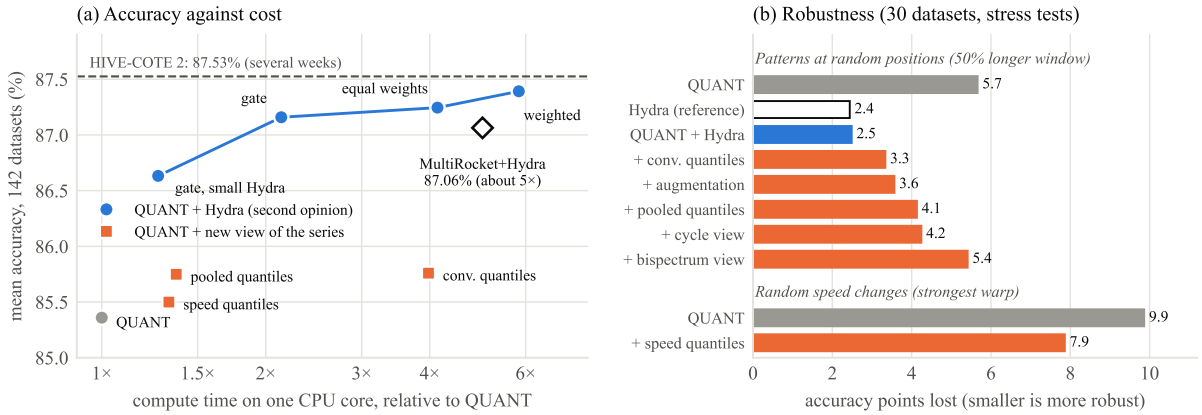


Fig. 2. (a) Mean accuracy against one-core compute time (both relative to plain QUANT in the same timing run). The blue line is QUANT + Hydra with more or less of Hydra used; orange squares add a new view of the series to QUANT. MultiRocket+Hydra's time is from the QUANT paper (other hardware). (b) Accuracy points lost in the two stress tests: patterns at random positions (window 50% longer) and series played at random speeds. Both panels cover ideas 1–5 and 14–18 only.

choices we made after seeing results were tested on unseen data (Sections V-B2, V-B6).

B. What Worked

Table III lists every idea and Fig. 2a shows accuracy against cost.

1) A second opinion from Hydra (idea 3): Inspiration: QUANT and Hydra fail on different data, and HC2 shows that averaging different methods pays off. *What:* we train both and average their class probabilities; Hydra's ridge scores become probabilities through a softmax whose "temperature" is fitted on held-out predictions, and each model is weighted by its cross-validated accuracy (CAWPE [13]). *Result:* 87.39%, +2.03 points (99/10/33, $p < 0.001$), +6.4 on the 18 many-class datasets (Pig +26 to +27), and statistically level with MultiRocket+Hydra ($+0.33$, $p = 0.37$) and HC2 (-0.13, $p = 0.63$) on the same splits. *Limitation:* 5.8× QUANT's time (91 instead of 16 minutes; 4.1× with equal weights, 87.25%): far cheaper than HC2, but no cheaper than MultiRocket+Hydra.

2) Asking Hydra only when QUANT is unsure (idea 18): Inspiration: the ensemble's gain is concentrated: looking back, 20 datasets gave 66% of it for 6% of Hydra's extra time. *What:* QUANT predicts first; only when the gap between its two most likely classes is below a threshold do we ask Hydra and average. As Hydra must process all training series whenever it is used, we also skip it for a whole dataset when too few of its test series are unsure (the *two-level gate*). Settings picked on one half of the datasets are applied to the other half. *Result:* 87.16%, +1.80 points (83/39/20, $p < 0.001$) at 2.1× QUANT's time: 95% of the ensemble's gain at half its cost, with Hydra asked about 26% of the test series. With its setting chosen on the first 60 datasets alone, the gate kept 97% of the gain on the other 82 (+2.19 points, $p < 0.001$), and a smaller Hydra (16 instead of 64 kernel groups) gives +1.27 points for 1.3×. *Conclusion:* level with MultiRocket+Hydra ($+0.10$, $p = 0.43$) at less than half its reported cost (other hardware), and 0.37 points behind HC2 ($p = 0.045$).

Do the gains stack? Adding the pooled and speed quantiles of ideas 1 and 14 (below) under the same gate gives 87.47%

TABLE III
ALL UNIVARIATE IDEAS: MEAN ACCURACY (%) ON 142 DATASETS, FIRST 10 RESAMPLES. Δ, W/D/L AND WILCOXON p ARE AGAINST QUANT; TIME IS ONE-CORE COMPUTE RELATIVE TO QUANT

Idea	Acc.	Δ	W/D/L	p	Time
QUANT (reference)	85.36				1×
Worked					
3: + Hydra, weighted	87.39	+2.03	99/10/33	<.001	5.8×
3: + Hydra, equal weights	87.25	+1.89	99/7/36	<.001	4.1×
18: + Hydra, two-level gate	87.16	+1.80	83/39/20	<.001	2.1×
18: + small Hydra, gate	86.63	+1.27	84/34/24	<.001	1.3×
1 + 14 + 18: stacked‖	87.47	+2.11	106/11/25	<.001	
1: + pooled quantiles	85.75	+0.39	84/19/39	<.001	1.4×
2: + convolution quantiles	85.76	+0.40	78/7/57	.004	4.0×
14: + speed quantiles	85.50	+0.14	81/9/52	.006	1.3×
9: keep top 25% of features	85.65	+0.29	86/15/41	<.001	0.6×¶
10: padded + native*	77.53	+1.97	10/0/1	.005	
same, 3 fresh datasets§	92.52	+0.21	2/0/1		
Did not help					
5: + augmentation	84.64	−0.72	46/9/87	<.001	5.8×†
5: + augm. + test-time avg.	85.01	−0.35	59/10/73	.03	5.9×†
10: native windows + length*	75.66	+0.11	5/0/6	.9	
15: + cycle view	85.36	0.00	61/16/65	.53	1.2×
16: + bispectrum view	85.47	+0.11	65/20/57	.24	1.2×
17: rank-uniform splits	85.27	−0.09	56/10/76	.053	1.2×
7: + autocorrelation, AR	85.61	+0.25	74/13/55	.059	
11: + multi-scale differences	85.44	+0.09	75/12/55	.15	
12: + order-aware counts	85.40	+0.04	63/15/64	.47	
13: Random Forest	84.09	−1.27	14/7/121	<.001	≈5×¶
13: LightGBM	82.66	−2.70	29/4/109	<.001	≈14×¶
MultiRocket+Hydra	87.06	+1.70			≈5×‡
HC2	87.53	+2.17			weeks‡

*The 11 variable-length datasets, 30 resamples (QUANT: 75.55%). §Three variable-length datasets not used before, 30 resamples (QUANT: 92.31%); too few for a test. †Measured on all cores. ‡From [3]; other hardware. ¶Training only, all cores, vs QUANT on the same computer (idea 9: in the same run). ‖Fixed gate (gap < 0.4, Hydra if ≥10% unsure), not picked on held-out datasets; idea 18 alone: 87.22%. Ideas 7, 11, 12 and stacking were not timed.

(+2.11 points, 106/11/25): +0.25 over idea 18 at the same setting (92/14/36, $p < 0.001$), never losing more than 0.5 points on a dataset. It is level with HC2 (-0.06 points, $p = 0.55$), but its cost was not measured.

3) Position-free features (ideas 1 and 2): Inspiration: ROCKET-style methods summarise each pattern detector over the whole series, which is why position does not matter to them. *What:* (a) *pooled quantiles*: slide windows of every size over the series and keep each quantile's maximum, minimum and mean over all positions ("somewhere there is a stretch like this"); (b) *convolution quantiles*: 288 MiniRocket kernels [11], each output summarised by nine quantiles. *Result:* (a) +0.39 points (84/19/39, $p < 0.001$) at 1.4×; (b) +0.40 points (78/7/57, $p = 0.004$) at 4.0×: the same gain for three times the cost.

4) Speed-invariant quantiles (idea 14): Inspiration: Proximity Forest [20], which compares series with time-warping distances, is 2.7 points worse than QUANT on average but beats it by more than a point on 23 datasets, mostly gestures and handwriting. *What:* extra window quantiles in which every value counts by how much the series moves there, not by how long it stays [22], so a gesture made twice as slowly gives the same features. *Result:* +0.14 points (81/9/52, $p = 0.006$) at 1.3×, +0.51 on the 23 datasets where the elastic method wins ($p = 0.002$), and 7.9 points lost instead of 9.9 in the speed test (Fig. 2b).

5) Feature pruning (idea 9): Training the forest takes most of QUANT's time and grows with the number of features, so a small forest (50 trees) ranks the features on the training data and QUANT's forest is trained on the top 50%, 25% or 10% only. Keeping 25% gives +0.29 points (86/15/41, $p < 0.001$) with 1.7× faster training (all cores, against plain QUANT re-timed in the same run); 10% is 2.0× faster (+0.18, $p = 0.15$).

6) Variable-length series: keep the padding, add native windows (idea 10): What: a padded series can hold 11 real points padded to 500, so we add windows that cover the same *fraction* of each series' own length next to the padded ones. *Result:* 77.53% against 75.55% on the 11 datasets × 30 resamples (+1.97 points, better on 10 of 11, $p = 0.005$). We designed this version after seeing that replacing the padding did not help (Section V-C), so we ran it unchanged on three unseen variable-length datasets (Asphalt road surfaces): there it gave only +0.21 points (clear only on AsphaltObstacles, +0.63, $p = 0.001$), although they are padded as heavily. *Conclusion:* worth adding for gesture-like data, where the same movement is made at different speeds; elsewhere it adds little.

C. Ideas That Did Not Work

Shift and stretch augmentation (idea 5). Four shifted or stretched (±10%) copies of every training series [23], optionally averaging over the same copies of each test series: -0.72 points ($p < 0.001$), or -0.35 with test-time averaging, at about 6× the cost. It helped where QUANT is weakest (PigAirwayPressure +33) but hurt where position *is* the meaning (Beef -17: a shifted spectrum is a different spectrum).

Removing the padding (idea 10, first version): +0.11 points with the series length as an extra feature ($p = 0.9$), -0.20 without, with big swings (GesturePebble +3.5, GestureMidAir -6.7): the padding tells the forest how long, and so how fast, a gesture was.

Cycle view (idea 15). QUANT trails HC2 most on series that repeat (-3.9 against -1.5 points elsewhere), so we folded each series into its median cycle, rotated to a standard start, as astronomers fold light curves [24]: ±0.00 (61/16/65), with gains on a few datasets (ToeSegmentation2 +5.2) and losses on outlines (Herring -2.8), but the most robust new view under random offsets (-4.2 against -5.7).

Shape-aware Fourier view (idea 16). The Fourier magnitudes cannot tell a wave from its mirror image; the bispectrum slice $X(f)^2\,\overline{X(2f)}$ ignores position but keeps shape [25]: +0.11 ($p = 0.24$), +1.0 on gestures but nothing on the many-class datasets it was meant for.

Rank-uniform splits (idea 17). ExtraTrees draws split points uniformly between a feature's extremes, which wastes splits on QUANT's skewed features; ranking every feature

TABLE IV
MULTIVARIATE DATA, MEAN ACCURACY (%). LEFT: 19 UEA DATASETS × 5 RESAMPLES (Δ, p AGAINST QUANT). RIGHT: THE 6 DATASETS WITH 3-AXIS SENSORS, EVERY SERIES TILTED AT RANDOM BY UP TO THE ANGLE SHOWN

	19 datasets			Sensors tilted by up to			
Method	Acc.	Δ	p	0°	90°	180°	Lost
QUANT	74.18			85.63	74.83	67.54	18.1
Hydra	74.03	−0.14	.73	87.89	76.99	69.12	18.8
18: + Hydra, gate	75.59	+1.41	.033	88.29	78.25	71.51	16.8
19: + cross-channel	76.80	+2.62	.031	85.92	78.22	73.59	**12.3**
18 + 19	**77.78**	+3.60	.001	88.27	79.66	**74.03**	14.2

first [26] gave -0.09 ($p = 0.053$): choosing the best of about 300 random splits already makes up for lopsided ones.

More views, other classifiers (ideas 7, 11, 12, 13). We added (a) the autocorrelation function and AR coefficients, which ignore where a repeating pattern starts, as in rSTSF [4]; (b) differences between moving averages over 2, 4 and 8 points, that is, medium-speed change located in time, like a Haar wavelet; (c) per-window counts that depend on the order of values (mean crossings, peaks, slope); and (d) replaced ExtraTrees by a Random Forest or LightGBM. (a) gave +0.25 ($p = 0.059$), though +14.3 on PigCVP and +4 to +6 on the other Pig recordings; (b) +0.09 ($p = 0.15$); (c) +0.04 ($p = 0.47$). (d) Random Forest lost 1.27 points (worse on 121 datasets) and LightGBM 2.70, collapsing with many classes and few examples (PigArtPressure -50, FiftyWords -35): ExtraTrees' random splits suit QUANT's many correlated features and small training sets.

D. What We Learned

First, QUANT's habit of looking at fixed positions is not simply a flaw: where a value occurs often *is* the class, so forcing QUANT to ignore position hurt, while adding a position-free view next to its own helped. Second, a second opinion beat every change to QUANT itself, and it is only needed where QUANT is unsure. Third, gains appear where the targeted weakness is. Fourth, robustness is not accuracy: every idea we stress-tested lost less than QUANT when patterns moved or sped up (Fig. 2b), but the cycle view and augmentation are more robust without being more accurate, as in most UCR datasets the patterns do not move much.

VI. MULTIVARIATE DATA

All datasets so far have one channel, but many sensors record several, such as the three axes of an accelerometer. QUANT's code computes its features on every channel separately and joins them, so it never sees how channels move together. We took 19 small and medium datasets from the UEA archive [27] (2–28 channels; not the largest or variable-length ones), each with its official split and 4 random splits with the same class balance. QUANT, the multivariate Hydra [9] and idea 18 run unchanged, with the gate settings from the univariate data.

Idea 18 carries over: +1.41 points (11/5/3, $p = 0.033$), although Hydra alone is no better than QUANT (Table IV). Hydra was asked about 48% of test series (26% on univariate data), so it saves less time.

Cross-channel views (idea 19). The same movement made with a tilted device changes every axis. We add new series built from several channels, described by QUANT's windows and quantiles: (a) for every 3-axis sensor, its magnitude $\sqrt{x^2 + y^2 + z^2}$ and the series along its own main direction of movement, which both stay the same when the sensor is rotated; (b) for all datasets, differences between pairs of channels (standardised per series; at most 15 pairs) and the series along the main direction of all channels. This gives +2.62 points (13/1/5, $p = 0.031$), more than idea 18, and the two stack: +3.60 together ($p = 0.001$; +2.19 over idea 18, $p = 0.032$). Most of the gain comes from (b) on datasets without 3-axis sensors (HandMovementDirection +23.8, NATOPS +8.5; +1.45 overall without HandMovementDirection); on the sensor datasets, already near their ceiling, it adds only +0.3.

Tilted sensors. To test (a), we rotated every series of the 6 sensor datasets at random by up to 30°, 90° or 180°, in training and test data. This changes QUANT's features by 36–69%, and those of the magnitude and main direction by 0%. At 180°, QUANT loses 18.1 points and idea 19 only 12.3, which puts it 6.1 points ahead, better on all 6 datasets ($p = 0.031$, the minimum with 6) and ahead of idea 18. Idea 19 still loses accuracy, because QUANT's own per-axis features stay in, and on Handwriting and UWaveGestureLibrary every method collapses, since there the direction of the movement is the class. *Limitations:* 19 datasets, our own splits, and half of the gain on normal data comes from one dataset.

VII. IDEAS WE COULD NOT TRY

A tabular foundation model such as TabPFN [28] or TabICL [29] on QUANT's features targets its worst case, many classes with few examples; RocketPFN [30] matched HC2 this way with ROCKET features, but at our dataset sizes these models need a GPU. The large MONSTER datasets [7] need more disk space than our laptop has.

VIII. CONCLUSION

QUANT's published results hold: 85.51% against the paper's 85.68%, with the gap coming from four corrected datasets. Its main weakness is that it looks at fixed positions, which hurts when patterns move and when there are many classes with few examples. Asking Hydra only when QUANT is unsure gives 87.16% (+1.8 points) at 2.1× QUANT's time, level with MultiRocket+Hydra at less than half its reported cost; adding pooled and speed quantiles reaches 87.47% (not timed), 0.06 points below HC2, which needs weeks. Keeping only a quarter of QUANT's features trains 1.7× faster and is slightly more accurate. Several intuitive ideas did not help. On multivariate data, views that combine channels gave +2.6 points (+3.6 with Hydra) and made QUANT far less sensitive to tilted sensors.

Work Plan and Contributions

Work Plan

Before starting, we met in person and planned the project together. The plan had three steps: reproduce QUANT, find out where it is weak, and only then try to improve it. We wanted each idea to fix a weakness we had measured, not one we had guessed.

The reproduction came first. We ran the authors' code on all 142 datasets and 30 resamples and compared our numbers with the paper's, dataset by dataset. The weakness analysis that followed decided what we worked on next, and it had two parts. We first went through QUANT's published results next to those of the best methods, grouped by type of data, number of classes and training set size. Gesture recordings, datasets with many classes but few examples, and padded datasets stood out. We then ran controlled tests that change one thing at a time, using Hydra as a reference because it does not care where in a series a pattern sits. We moved patterns to random positions, cut training sets down to a few examples per class, compared lightly and heavily padded series, removed test series that also appear in the training set, and normalised the series. The clearest finding: QUANT depends on where things happen in a series.

Our ideas grew out of these results. Every promising one went into an `ideas.md` file, with the weakness it targets and the papers behind it, and we read those papers to judge whether an idea could work. To fit our availability we split the research into two parts, and we met every week to go through new results and decide what to do next. We also set ourselves some rules, mostly so we would not fool ourselves. Every test uses the same splits and seeds as the reproduction and a paired statistical test, failures are reported like successes, and any choice made after seeing results is checked again on data it has not seen.

Work Division

We both agree that the work was divided fairly and equally, 50% each. Both of us did research, coding and evaluation, and we planned the final report together. The research and the conclusions were shared, while the coding and the experiments were divided by time, which suited our timetables. Dániel had more free time early in the semester, so he started with the reproduction, the weakness analysis and the first two rounds of ideas. He left out ideas our hardware could not handle, such as pairing QUANT with a tabular foundation model, which needs a GPU. Tymur had more time later on. He built on Dániel's results, picked up ideas that were still open and extended our tests to multivariate data. He also brought ideas from later lectures into the project and added idea 19 to `ideas.md`. We went through every result together, since we both wanted to learn from the whole project and not only from our own half. Our `ideas.md` lists every idea we discussed, including the ones we never implemented.

Individual Contributions

Dániel Sajben built the reproduction pipeline and the weakness tests and started `ideas.md`. In two rounds he implemented and evaluated ten ideas (1, 2, 3, 5 and 10, then 14–18). Idea 18 worked best: asking Hydra only when QUANT is unsure gave 87.16% at 2.1× QUANT's compute time. Position-free and speed-invariant features added a little, and four ideas did not help. He wrote the first draft of this report and set up the repository so that every number can be checked or rerun.

Tymur Sukhariev ran rounds 3 and 4, which tested the ideas that were still open. On the 142 univariate datasets he implemented and evaluated ideas 7, 11, 12 and 13 (more views and other classifiers, none of which helped), idea 9, which keeps a quarter of QUANT's features and trains 1.7× faster with slightly higher accuracy, and the stacked ideas 1 + 14 + 18, our best result (87.47%). He took our tests to multivariate data: he set up 19 datasets from the UEA archive [27], showed that our gated ensemble (idea 18) works there too (+1.4 points), and proposed and tested idea 19, views that combine channels (+2.6 points, +3.6 with Hydra), including a stress test with tilted sensors. He also made our setup run on macOS and merged the new results into this report.

We both agree on an equal split: Dániel Sajben 50%, Tymur Sukhariev 50%.

REFERENCES

[1] M. Middlehurst, J. Large, M. Flynn, J. Lines, A. Bostrom, and A. Bagnall, "HIVE-COTE 2.0: A new meta ensemble for time series classification," *Machine Learning*, vol. 110, pp. 3211–3243, 2021.

[2] M. Middlehurst, P. Schäfer, and A. Bagnall, "Bake off redux: A review and experimental evaluation of recent time series classification algorithms," *Data Mining and Knowledge Discovery*, vol. 38, pp. 1958–2031, 2024.

[3] A. Dempster, D. F. Schmidt, and G. I. Webb, "QUANT: A minimalist interval method for time series classification," *Data Mining and Knowledge Discovery*, vol. 38, pp. 2377–2402, 2024.

[4] N. Cabello, E. Naghizade, J. Qi, and L. Kulik, "Fast, accurate and explainable time series classification through randomization," *Data Mining and Knowledge Discovery*, vol. 38, pp. 748–811, 2024.

[5] P. Geurts, D. Ernst, and L. Wehenkel, "Extremely randomized trees," *Machine Learning*, vol. 63, no. 1, pp. 3–42, 2006.

[6] M. Middlehurst, A. Ismail-Fawaz, A. Guillaume, C. Holder, D. Guijo-Rubio, G. Bulatova, L. Tsaprounis, L. Mentel, M. Walter, P. Schäfer, and A. Bagnall, "aeon: A Python toolkit for learning from time series," *Journal of Machine Learning Research*, vol. 25, no. 289, pp. 1–10, 2024.

[7] A. Dempster, N. M. Foumani, C. W. Tan, L. Miller, A. Mishra, M. Salehi, C. Pelletier, D. F. Schmidt, and G. I. Webb, "MONSTER: Monash scalable time series evaluation repository," *Journal of Data-centric Machine Learning Research*, vol. 2, pp. 1–47, 2025.

[8] H. A. Dau, A. Bagnall, K. Kamgar, C.-C. M. Yeh, Y. Zhu, S. Gharghabi, C. A. Ratanamahatana, and E. Keogh, "The UCR time series archive," *IEEE/CAA Journal of Automatica Sinica*, vol. 6, no. 6, pp. 1293–1305, 2019.

[9] A. Dempster, D. F. Schmidt, and G. I. Webb, "Hydra: Competing convolutional kernels for fast and accurate time series classification," *Data Mining and Knowledge Discovery*, vol. 37, pp. 1779–1805, 2023.

[10] A. Dempster, F. Petitjean, and G. I. Webb, "ROCKET: Exceptionally fast and accurate time series classification using random convolutional kernels," *Data Mining and Knowledge Discovery*, vol. 34, pp. 1454–1495, 2020.

[11] A. Dempster, D. F. Schmidt, and G. I. Webb, "MiniRocket: A very fast (almost) deterministic transform for time series classification," in *Proc. 27th ACM SIGKDD Conf. Knowledge Discovery and Data Mining*, 2021, pp. 248–257.

[12] C. W. Tan, A. Dempster, C. Bergmeir, and G. I. Webb, "MultiRocket: Multiple pooling operators and transformations for fast and effective time series classification," *Data Mining and Knowledge Discovery*, vol. 36, pp. 1623–1646, 2022.

[13] J. Large, J. Lines, and A. Bagnall, "A probabilistic classifier ensemble weighting scheme based on cross-validated accuracy estimates," *Data Mining and Knowledge Discovery*, vol. 33, pp. 1674–1709, 2019.

[14] A. Dempster, "QUANT source code," https://github.com/angus924/quant, 2024.

[15] A. Ismail-Fawaz, A. Dempster, C. W. Tan, M. Herrmann, L. Miller, D. F. Schmidt, S. Berretti, J. Weber, M. Devanne, G. Forestier, and G. I. Webb, "An approach to multiple comparison benchmark evaluations that is stable under manipulation of the comparate set," arXiv:2305.11921, 2023.

[16] A. Guillaume, C. Vrain, and W. Elloumi, "Random dilated shapelet transform: A new approach for time series shapelets," in *Proc. Int. Conf. Pattern Recognition and Artificial Intelligence (ICPRAI)*, LNCS 13363, 2022, pp. 653–664.

[17] P. Schäfer and U. Leser, "WEASEL 2.0: A random dilated dictionary transform for fast, accurate and memory constrained time series classification," *Machine Learning*, vol. 112, pp. 4763–4788, 2023.

[18] H. Ismail Fawaz, B. Lucas, G. Forestier, C. Pelletier, D. F. Schmidt, J. Weber, G. I. Webb, L. Idoumghar, P.-A. Muller, and F. Petitjean, "InceptionTime: Finding AlexNet for time series classification," *Data Mining and Knowledge Discovery*, vol. 34, pp. 1936–1962, 2020.

[19] M. Middlehurst and A. Bagnall, "The FreshPRINCE: A simple transformation based pipeline time series classifier," in *Proc. Int. Conf. Pattern Recognition and Artificial Intelligence (ICPRAI)*, LNCS 13364, 2022, pp. 150–161.

[20] B. Lucas, A. Shifaz, C. Pelletier, L. O'Neill, N. Zaidi, B. Goethals, F. Petitjean, and G. I. Webb, "Proximity Forest: An effective and scalable distance-based classifier for time series," *Data Mining and Knowledge Discovery*, vol. 33, pp. 607–635, 2019.

[21] Y. Zhang, G. Batista, and S. S. Kanhere, "Revisit time series classification benchmark: The impact of temporal information for classification," in *Proc. Pacific-Asia Conf. Knowledge Discovery and Data Mining (PAKDD)*, 2025, pp. 193–205.

[22] S. Banach, "Sur les lignes rectifiables et les surfaces dont l'aire est finie," *Fundamenta Mathematicae*, vol. 7, pp. 225–236, 1925.

[23] A. Le Guennec, S. Malinowski, and R. Tavenard, "Data augmentation for time series classification using convolutional neural networks," in *ECML/PKDD Workshop on Advanced Analytics and Learning on Temporal Data*, 2016.

[24] U. Rebbapragada, P. Protopapas, C. E. Brodley, and C. Alcock, "Finding anomalous periodic time series," *Machine Learning*, vol. 74, pp. 281–313, 2009.

[25] P. A. Parker, S. H. Holan, and N. Ravishanker, "Nonlinear time series classification using bispectrum-based deep convolutional neural networks," arXiv:2003.02353, 2020.

[26] T. Galili and I. Meilijson, "Splitting matters: How monotone transformation of predictor variables may improve the predictions of decision tree models," arXiv:1611.04561, 2016.

[27] A. Bagnall, H. A. Dau, J. Lines, M. Flynn, J. Large, A. Bostrom, P. Southam, and E. Keogh, "The UEA multivariate time series classification archive, 2018," arXiv:1811.00075, 2018.

[28] N. Hollmann, S. Müller, L. Purucker, A. Krishnakumar, M. Körfer, S. B. Hoo, R. T. Schirrmeister, and F. Hutter, "Accurate predictions on small data with a tabular foundation model," *Nature*, vol. 637, pp. 319–326, 2025.

[29] J. Qu, D. Holzmüller, G. Varoquaux, and M. Le Morvan, "TabICL: A tabular foundation model for in-context learning on large data," in *Proc. 42nd Int. Conf. Machine Learning (ICML)*, PMLR 267, 2025, pp. 50817–50847.

[30] F. M. O'Rourke, A. Trisovic, and D. Bertsimas, "RocketPFN: Accurate time series classification via in-context learning," arXiv:2606.21786, 2026.